

# Project 1: Deep-Q-Learning-for-Navigation

## Udacity Deep Reinforcement Learning Nanodegree Program

Bob Flagg \*




Figure 1: Unity ML-Agent: Banana Collector

**Keywords:** Reinforcement Learning, Deep Learning

### 1 Introduction

In this project I'll present the following three solutions to the Unity ML-Agent Banana Collector environment:

1. *Deep Q-Learning* [Mnih et al. 2015],
2. *Double Deep Q-Learning* [van Hasselt et al. 2015], and
3. *Dueling Deep Q-Learning* [Wang et al. 2015].

Source code in Python, using PyTorch, is available on github in the repo Deep-Q-Learning-for-Navigation.

### 2 Background

The Unity ML-Agent Banana Collector is a *sequential decision making problem*, in which an agent interacts with an environment over discrete time steps and seeks to maximize the expected *discounted return*:

$$G_t = \sum_{\tau=t}^{\infty} \gamma^{\tau-t} R_{\tau},$$

where  $\gamma \in [0, 1]$  is a discount factor that trades-off the importance of immediate and future rewards. See [Sutton and Barto 1998] for a general discussion of this sort of problem. In this specific example, the agent observes a 37 dimensional vector containing the agent's velocity, along with ray-based perception of objects around the agent's forward direction and tries to learn how to best select one of the following actions:

1. move forward,
2. move backward,
3. turn left, and
4. turn right.

A reward of +1 is provided for collecting a yellow banana, and a reward of -1 is provided for collecting a blue banana. Thus, the

goal of the agent is to collect as many yellow bananas as possible while avoiding blue bananas.

### 3 Deep Q-Learning for Navigation




Figure 2: Deep Q-Network [Wang et al. 2015]

In deep Q-learning the action-value function,  $Q(s, a)$  is approximated with a *deep Q-network*,  $Q(s, a|\theta)$ , with parameters  $\theta$ . To train this *local network*, we optimize the following sequence of loss functions at iteration  $i$ :

$$L_i(\theta_i) = \mathbb{E}[(y_i^{DQN} - Q(s, a|\theta_i))^2],$$

with

$$y_i^{DQN} = R_i + \gamma \cdot \max_{a'} Q(s', a'|\theta_i^-),$$

where  $\theta^-$  represents the parameters of a fixed and separate *target network*.

\*Deep-Q-Learning-for-Navigation




Figure 3: Scores for Deep Q-Learning

## 4 Double Deep Q-Learning for Navigation

In Q-learning and deep Q-learning, the max operator uses the same values to both select and evaluate an action. This can lead to over optimistic value estimates. *Double deep Q-learning* mitigates this problem by using a different target:

$$y_i^{DDQN} = R_i + \gamma \cdot Q(s', \arg \max_{a'} Q(s', a' | \theta_i) | \theta^-).$$




Figure 4: Scores for Double Deep Q-Learning

## 5 Dueling Deep Q-Learning for Navigation




Figure 5: Dueling Deep Q-Network [Wang et al. 2015]

The architecture of the Q-network in dueling deep Q-learning has two streams, one estimates the state-value function,  $V(s|\theta, \beta)$ , and the other estimates an *advantage function*, which in one implementation has the form

$$A(s, a|\theta, \alpha) = \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'|\theta, \alpha).$$

The estimate of the action-value function is then

$$Q(s, a|\theta, \alpha, \beta) = V(s|\theta, \beta) - (A(s, a|\theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'|\theta, \alpha)).$$

**Note:** I don't use convolution in the network for the Unity ML-Agent Banana Collector so there are no shared  $\theta$  parameters in this example.




Figure 6: Scores for Dueling Double Deep Q-Learning

## 6 Improving Performance

### References

- MNIH, V., KAVUKCUOGLU, K., SILVER, D., RUSU, A. A., VENESS, J., BELLEMARE, M. G., GRAVES, A., RIEDMILLER, M. A., FIDJELAND, A., OSTROVSKI, G., PETERSEN, S., BEATTIE, C., SADIK, A., ANTONOGLU, I., KING, H., KUMARAN, D., WIERSTRA, D., LEGG, S., AND HASSABIS, D. 2015. Human-level control through deep reinforcement learning. *Nature* 518, 7540, 529–533.
- SUTTON, R. S., AND BARTO, A. G. 1998. *Reinforcement learning - an introduction*. Adaptive computation and machine learning. MIT Press.
- VAN HASSELT, H., GUEZ, A., AND SILVER, D. 2015. Deep reinforcement learning with double q-learning. *CoRR* abs/1509.06461.
- WANG, Z., DE FREITAS, N., AND LANCTOT, M. 2015. Dueling network architectures for deep reinforcement learning. *CoRR* abs/1511.06581.